

The Executive View — product narrative and screen-by-screen design brief

For: a Claude Design session that will draw every screen in this document. **Audience of the final product:** McKinsey's Leader of Tech in Healthcare and Leader of Product in Healthcare. **Hard constraint:** the code is frozen. Nothing new gets built. Every change in this document is layout, naming, grouping, type, colour and spacing over data the system already returns.

How to read this document

- **Part 1** is the product narrative. Read it first. Every design decision is downstream of it.
- **Part 2** is the design law — the few things that are fixed. Light scheme, locked type and spacing, and the hard constraints. Everything not named there is yours to decide.
- **Part 3** is the screens. Four screens, each broken into its sections. Every section has a screenshot, what it holds today, and **who reads it and what they need from it**. The visual answer is yours.
- **Part 4** is the plain-English dictionary. The single largest change in this redesign is replacing codes with words. This is the word list.

How to treat Part 3. Where it describes today's screen, that is fact — read it as the inventory. Where it suggests a treatment, read it as *intent, not instruction*: it tells you what the section is for and what has to survive. If you find a better way to serve the same purpose, take it. The only lines that are binding are the ones marked as constraints — available data fields, and the rules in Part 2.

Part 1 — The product narrative

What a product narrative is

A product narrative is not a feature list and not a pitch deck. It is the **story of a person's problem, told in their language, in which your product is the turn**. It has a fixed shape, and every good one has all five parts:

1. **The world, as the buyer already experiences it.** No product in it yet. If they don't nod here, nothing after lands.
2. **The tension.** The specific thing that is broken or expensive, stated as a cost they already pay.
3. **The turn.** What changes. One sentence. Not a feature — a change in what is possible.
4. **The new world.** What their day looks like after. Concrete, not adjectival.
5. **The proof.** Why this is real and not a demo.

Why the CTO told you to ask for one: a narrative forces you to name the buyer, the cost and the change *before* you design anything. Once it exists, it settles arguments. "Does this screen belong?" becomes "does this screen serve the narrative?" — which has an answer. Every section of Part 3 is justified against Part 1.

The other reason: your audience is a Tech leader and a Product leader at a consulting firm. They are not grading your engineering. They are grading whether you can **see a business through a system**. An engineering narrative says "here is what I built." A product narrative says "here is what it is worth, to whom, and why the alternative is expensive."

The narrative

1. The world

A hospital pharmacy hands an expensive drug to a patient. That is one physical act, and it takes two seconds.

From that moment the hospital is owed money — and the money comes back from **up to three different places, on three different timetables, in three different formats, from parties who do not talk to each other**.

- Pills picked up at a counter bill to a **PBM**.
- The same drug infused in a chair bills to a **medical insurer**, through a clearinghouse.
- On top of either, if the hospital qualifies for the federal **340B** programme, a **drug manufacturer** owes a separate cash rebate — and whether this particular dispense qualified is ruled on by a **TPA**, a company whose entire job is that one decision.
- A fourth feed, the **bank**, owes nothing. It is the only proof any of the other three actually landed.

Roughly 30–60% of a health-system specialty pharmacy's claims carry a 340B rebate. So for a large share of the book, the hospital is chasing two receivables from two unrelated counterparties on one drug.

2. The tension

Nobody sends the hospital a statement saying: here is what you were owed, and here is what you got.

There is no such document, because there is no shared identifier. The four feeds were standardised by different bodies, for different rails, in different decades. The pharmacy feed and the medical feed share literally nothing. A bank line does not carry a claim number at all — it carries a trace number that resolves to a *remittance*, and the remittance holds the claim list. One deposit can cover dozens of claims.

So the month-close question — *we dispensed millions of dollars of drug; where is it?* — is answered today with a spreadsheet and a `VLOOKUP`. And that spreadsheet cannot distinguish:

- **an underpayment from a contractual discount.** A payer that pays less than you billed is usually correct — the contract says so. The test is arithmetic: charge minus payment minus stated adjustments. A non-zero remainder is money genuinely missing. Nobody does that subtraction on every claim by hand.
- **a rebate that was refused from a rebate that was approved and never sent.** The manufacturer's paperwork says paid. No cash ever arrives. Nothing on the reimbursement side would ever surface it, because it is a different counterparty on different paperwork.
- **a missing claim from a failed match.** The money arrived. The identifier was spelled `07845102` in one feed and `7845102` in the other. The claim still reads as *awaiting payment* forever. The claim is not missing — the mapping failed, and that is a different problem with a different owner.
- **a clawback from forty small underpayments.** A payer that thinks it overpaid you last month does not send an invoice. It **pays you less in a later batch** and explains the difference in a line that belongs to no claim in that batch. Reconcile the deposit against the claim list without reading that line, and one traceable recoupment reads as forty phantom underpayments.
- **a compliance problem from a good month.** The insurer refused to pay for the drug. The manufacturer sent the 340B rebate anyway. Net position on that dispense is negative — the hospital is holding money it may owe back, on a claim that paid nothing. No single-track view can see it.

Each of those is money. The last one is money *and* an audit finding.

3. The turn

The statement nobody sends is the product.

One system reads all four feeds, rebuilds each dispense into a single financial story despite the missing identifier, and states — per claim, in dollars — what should have arrived, what did, what is still outstanding, and which of those needs a human to pick up the phone.

Three of those four are arithmetic. **Only the last one is a judgement call.** That split is the whole design.

4. The new world

The analyst opens one screen and sees the book sorted into three piles, because "*what do I do with this?*" has exactly three answers: **nothing, wait, work it.**

They open the pile that needs work, ranked by money. They click a claim and get its entire life in one view — what was filed, what the payer said, when cash moved, how the rebate went, and every point at which the answer changed. They can drag the date backwards and see exactly what the system believed on 31 March, because that is a parameter, not a feature.

When they want the reasoning in English, an AI layer reads that same evidence and writes four paragraphs with citations that click through to the literal line in the source file. When they want a recommendation, a second model grades the first against a sixteen-point checklist before a human ever sees it. Then a **person** — editing every field, not approving a draft — commits the one row the AI layer is ever allowed to write.

5. The proof

Three claims, each one measurable rather than assertable:

- **Every number is computed in Python. The model only explains numbers it was handed.** Nine tools; eight read-only; zero write tools in any model's tool list. A Python check scans every figure in generated prose against what the tools actually returned, and an unsourced figure is a veto, not a deduction. *A number a model produced is a number you cannot defend in an audit.*
- **Correctness is measured, not claimed.** The rules were enumerated before they were written: 12,093,235,200 raw combinations reduce to 4,224 legal ones, producing 372 verdict pairs — all 372 reachable. The engine reproduces the independent oracle on 4,224 of 4,224, with no tolerance. One disagreement fails the build.

- **The system never guesses, and says so.** Two claims that are genuinely indistinguishable are parked, not picked — because picking one attributes real money to the wrong claim and then looks identical to a correct match in every report afterwards. *"I cannot tell"* is a first-class answer with its own code, not a failure state.

The line to end on

The deterministic layer decides what is true. The AI layer decides what is worth saying about it. A human decides what to do.

What this narrative demands of the interface

Every design instruction in Part 3 comes from one of these. If a change cannot be traced to a line here, it is decoration and should be cut.

THE NARRATIVE SAYS	SO THE SCREEN MUST
"what should have arrived, what did, what's outstanding"	lead with money , not counts. Dollars first, episodes second.
"nothing, wait, work it"	make the three piles the most prominent object on the page, and name them as actions, not as database enums.
"which of those needs a human"	rank by money and age, and say out loud what the human is supposed to do.
"a different problem with a different owner"	keep the failure types visibly distinct. Never collapse them into "unmatched".
"what did we believe on 31 March"	make the date control read as a point in time, not as a slider with a number on it.
"the model only explains numbers it was handed"	keep the provenance visible — citations, source links, "computed in Python" — because that is the differentiator, not a footnote.
"an audit finding"	compliance flags cannot look like ordinary chips.

Part 2 — Design law

This is not a greenfield visual system. It was measured.

`docs/frontend_linear_benchmark.md` records a real browser measurement of linear.app — 3,113 elements across three passes — diffed against this app. **Do not re-derive these values and do not "improve" them.** They are already correct and already implemented.

Locked — use as-is

THING	VALUE
UI typeface	Inter Variable (local npm package, already installed)
Identifier typeface	<code>ui-monospace</code> stack. Reserved for ids, dates and money — the role Linear gives Berkeley Mono
Weights	400 body · 550 buttons · 640 headings
Type scale	10 / 11 / 12 / 13 / 14 / 16 / 19 / 28px. Eight steps, no ninth
Letter-spacing	−0.01em body · −0.014em at 19px · −0.022em reserved for 28px display only
Line height	1.25 tight · 1.55 body
Spacing scale	2 / 4 / 6 / 8 / 12 / 16 / 24 / 32 / 48px
Radius	12px cards and panels · 8px md · 6px controls · 4px xs · 999px pills
Page background	<code>#f7f8fa</code>
Ink	<code>#16161c</code>
Shadows	Zero. A deliberate, reasoned divergence — a shadow reads as depth on a near-black page and as smudge on a light one at this density. Borders do the separating.
Theme	Light only, unconditionally. Dark mode was removed on purpose. Do not add a dark variant.

The full colour set, and what each is for

TOKEN	HEX	USED FOR
--surface-0	#f7f8fa	page background
--surface-1	#ffffff	panels, cards, buttons, tiles
--surface-2	#f0f1f4	nested content — table heads, chips, code pills, hover, log blocks
--border	#e4e5ea	every hairline
--border-strong	#c6c8d2	control borders, and the fallback for every runtime custom property
--text-primary	#16161c	body ink
--text-secondary	#4c4d5a	headings, table headers, tile labels
--text-muted	#66677a	hints, empty states, timestamps (4.86:1 on --surface-2)
--status-good	#0ca30c	Closed
--status-warning	#fab219	Pending
--status-serious	#ec835a	cross-track flag chips
--status-critical	#d03b3b	Exception, reopened, negative variance, errors
--track-rebate	#4a3aa7	the 340B rebate track — declared and currently unused by the stylesheet
--accent	#5e6ad2	links, primary buttons, selection, focus
--accent-hover	#4f5ac0	primary hover
--accent-soft	#ecee fb	selected table row wash

Four more are set at runtime from JSX and consumed in CSS with fallbacks: --tile-color , --dot , --d , --oc . Any recolour must audit those too.

What you are allowed to change

Everything about **arrangement, naming, hierarchy and grouping**. Nothing about the token values above.

In scope: which section sits where; what is a heading versus a label; what is a chip versus plain text; what is shown by default versus behind a disclosure; column order, count and names; what is large and what is small; what is --text-primary versus --text-muted ; how many words a label gets; what gets a colour and what stays neutral.

What you may not do

1. **No new data.** The backend is frozen. Every field you can use is listed in Part 3 under its section. If a field is not listed there, it does not exist and cannot be fetched.
2. **No new dependencies.** The entire dependency list is react , react-dom , @fontsource-variable/inter . There is **no charting library, no icon library, no CSS framework, no router, no state manager**. Every icon on screen today is a literal Unicode character (▲ ☑ ✓ ◀ ▶ ⚠ ✕ → ← ⏪ ⏩). Every bar is hand-rolled CSS. A chart, if one is needed, must be inline SVG or CSS.
3. **No renaming of verdict, reason or exception codes.** The code strings feed the AI layer's grounding tables and recorded evaluation fixtures. Renaming one silently breaks the AI layer. **Add plain-English labels beside the codes; never replace the code string in the data.**
4. **No dark mode.**
5. **No shadows.**

The three aesthetic rules that carry the whole redesign

- 1. Money is the headline. Everything else is a subtitle.** Today the biggest number on a queue tile is an episode count (40) and the dollar figure is 12px monospace underneath it. Backwards for this audience. The dollars are why anyone opened the page.
 - 2. Every code gets a word. The code stays, demoted.** The screen is currently readable only by someone holding a 43-entry lookup table. The plain-English strings already exist in the codebase — `VERDICT_DESCRIPTIONS` at `src/recon/domain/verdicts.py:126`, exposed by `describe()` at line 142. The frontend has never imported them. Part 4 is the full word list. Pattern: **the human label is the primary text; the code becomes a small muted monospace suffix** so an operator can still cite it.
 - 3. One thing per screen region, and a stated question above it.** Each panel today opens with a 3–5 line paragraph of design rationale — genuinely good writing, aimed at a reviewer grading the architecture. An executive does not read it. Convert each to **one short line stating the question the panel answers**; move the rationale to a hover or drop it.
-

Part 3 — The screens

Four screens after this redesign, up from three. Splitting Verdict distribution onto its own screen is the one change here that touches wiring rather than pixels; its cost is stated in Screen 3.


#	SCREEN	ROUTE	THE QUESTION IT ANSWERS
1	Operations	/	Where is my money, and what needs working?
2	Episode	/analyse/{id}	What happened to this one claim, and what should we do?
3	Verdict library	/ + nav state (new)	What kinds of problems does the book contain?
4	Connectivity	/ + nav state	Which data sources are actually wired up?

Screens 1, 3 and 4 share one masthead and swap the body below it. Screen 2 is a separate page reached by a button. All figures quoted below are live, read off the running app on the demo profile, seed 8984883 , cursor 2026-07-01 .

Screen 1 — Operations (the landing page)

[illegible]

Today: one 4,727px scroll containing seven stacked panels in a 3fr/2fr grid, with a sticky episode panel on the right.



Operations, what you actually see first

The single biggest structural problem: above the fold you get the masthead, a paragraph about how the replay cursor works, and a paragraph about how the Portfolio Analyst works. The three tiles carrying the money are below the fold. An executive has to scroll to reach a dollar figure.


The fix in one sentence: the money tiles move to the top directly under the masthead; the replay cursor becomes a compact control inside the masthead row; the Portfolio Analyst drops below the worklist.

Target order for the left column

- 1. Money summary — the three tiles, enlarged
- 2. The worklist — queue table
- 3. Money and documents we could not match — renamed feed exceptions
- 4. Portfolio Analyst
- 5. To-do list

Verdict distribution leaves this screen entirely.

1.1 Masthead



Masthead

Today. `Post-Claim Pharmacy Financial Reconciliation` at 19px/640. Below it: "Deterministic layer. Every figure below was computed in Python and is shown verbatim — nothing on this screen calculates. Profile `demo`, seed `8984883`, engine

1.0.0." Right side, baseline-aligned: an Operations/Connectivity segmented toggle, the text 60 episodes · \$627,524.26 total variance, and three buttons — Rebuild demo, Rebuild full, Repeat seed.

Why it fails.


- The subtitle is an engineering disclaimer plus three build parameters. seed 8984883 means nothing to a buyer and signals "test harness".
- The only real money figure on the masthead — \$627,524.26 — renders as 12px muted body text, smaller and quieter than the buttons beside it.
- Three destructive Rebuild buttons occupy the most valuable real estate on the page. One wipes and regenerates the entire database.
- Page <h1> and browser tab title disagree — "Financial" is in one and not the other.

Build this instead.

- Product name left, one line. Drop "Deterministic layer" and the build parameters from the visible subtitle — put profile/seed/engine behind a small muted disclosure. Keep the "every figure was computed in Python" claim; it is a differentiator. Make it one quiet line, not the first thing under the title.
- Promote \$627,524.26 to a display figure — 28px, tabular, --text-primary, with Total variance across 60 claims beneath it at 11px uppercase muted. This is the number the room should see first.
- Move the replay cursor into this row (\$1.2).
- Collapse the three Rebuild buttons into one secondary control, or move them off the masthead entirely. They are demo plumbing, not product.
- Fix the tab title to match the <h1>.
- The masthead row is align-items: baseline while the control cluster inside it is align-items: center. Pick one.

Available fields. GET /api/meta → profile, episode_count, engine_version, adapter_version, stored.master_seed, cursor.{min, max, window_start, window_end}. Totals from GET /api/overview.

1.2 Replay cursor



Replay cursor

Today. A full panel with <h2>Replay cursor and a three-line paragraph. The control: the word Cursor, a monospace 2026-07-01, a ◀ button, a range slider, a ▶ button, and beneath it two bare dates 2025-07-01 and 2026-07-01 at the far edges with nothing labelling them.

Why it fails.

- A whole panel — ~195px above the fold — for one date control.
- "Cursor" is a database word. Three names for one control: the heading says Replay cursor, the visible label says Cursor, the accessible name says Replay cursor, and the code comment says there is no replay mode.
- The two end dates read as orphans. Nothing says they are the window bounds.
- Bare ISO dates. 2026-07-01 is not how a person says a date.
- The slider is not disabled while data loads, though the step buttons are — dragging mid-fetch is possible.

Build this instead.

- Move it into the masthead row as a compact control. No panel, no heading.
- Label it as a question: `Showing the book as of` followed by the date.
- Format the date long — `1 July 2026` — keeping tabular figures so it does not jitter while dragging.
- Add a quiet persistent state when the cursor is not at the latest date: a chip reading `Historical view – 142 days back` in `--status-warning`. Show nothing at the newest date. This is the feature that proves the audit story, and right now it is invisible unless you already know what the slider does.
- Label the end dates inline: `Earliest 1 Jul 2025 ... Latest 1 Jul 2026`.
- Disable the slider along with the step buttons while busy.

Available fields. `cursor.window_start`, `cursor.window_end`, `cursor.min`, `cursor.max` from `/api/meta`. Format is `YYYY-MM-DDTHH:MM:SSZ`; the control always emits end-of-day. Every read endpoint accepts it except `/trace`, `/record`, `/connectivity` and `/meta`.

1.3 The three money tiles — the most important object on the page

oldest 40 episodes

DATE	AGE	REIMB.	REBATE	VARIANCE	WHY
08-09	326d	A-17	C-01	-\$9,087.06	reopened from CLOSED DUPLICATE_PAYMENT REBATE_AWAITING_QUALIFICATION
10-02	272d	B-15	C-00	\$0.00	reopened from CLOSED CASH_MISMATCH
07-04	362d	B-09	C-00	\$2,203.20	reopened from PENDING APPEAL_LOST
07-06	360d	A-09	C-13	-\$1,197.56	reopened from PENDING OVERPAID REBATE_CLAWED_BACK
07-11	355d	A-04	C-09	\$4,149.00	reopened from PENDING REBATE_NO_CASH
07-31	335d	A-15	C-00	\$6,706.35	reopened from PENDING ADJUSTMENT_RESIDUAL

Queue tiles

Today. Heading `Queues at this cursor`, a three-line paragraph, then three tiles in a `repeat(auto-fit, minmax(230px, 1fr))` grid. Each tile: a 12px uppercase label with a Unicode glyph, a 28px count, a 12px monospace money line.

TILE	GLYPH	COUNT	MONEY LINE	COLOUR
Exception	▲	40	\$517,240.34 variance · 31 reopened	--status-critical
Pending	⌚	10	\$84,590.23 variance · 5 reopened	--status-warning
Closed	✓	10	\$25,693.69 variance	--status-good

Why it fails.

- The count is 28px and the dollars are 12px. The eye lands on `40`, `10`, `10`. The narrative says money leads.
- `Exception`, `Pending`, `Closed` are database enum names describing a *state*. The narrative promised an *action*: nothing, wait, work it.
- The only explanation of each bucket is a `title` tooltip — invisible until hover, unreachable on touch, and it is the most valuable copy on the screen. The three strings: "A defect exists. Work it." / "Waiting on an external party. No defect. Ranked by age." / "Nothing to do."
- `variance` is jargon. So is `reopened` with no explanation.
- Clicking the already-selected tile re-fetches rather than doing nothing.

Build this instead.

- **Invert the hierarchy.** Dollars at 28px tabular. Episode count demoted to a 12px line: `40 claims`.
- **Rename to the action, keep the state as a subtitle.** Suggested: `Needs work` / `Waiting on someone` / `Settled`, with the enum name as a small muted suffix.
- **Promote the tooltip to visible copy** — one line under each tile at 11px muted. This single change does more for comprehension than anything else on the screen.
- **reopened needs words.** It means *this claim was settled and has come undone*. 31 of 40 exceptions are reopened — the loudest signal on the page, currently a six-word monospace fragment. Give it a line: `31 were previously settled and reopened`.
- Keep the 3px coloured left border and the colour assignment exactly as-is. They work.
- Make the selected tile obvious beyond the current `--surface-2` fill — selection currently reads as hover.

Available fields. `GET /api/overview` → `by_disposition.{EXCEPTION|PENDING|CLOSED}.{episodes, expected_cents, received_cents, variance_cents, reopened}`. All three keys always exist, zero-filled.

Two cautions. `variance_cents` here is a sum of **absolute** values; the per-row `total_variance_cents` on a queue row is **signed**. They will not reconcile — never present them as the same quantity. And `expected_cents` / `received_cents` are already returned per bucket and **displayed nowhere today** — they are free material for an "\$X expected, \$Y received" line at zero backend cost.

1.4 The worklist (queue table)

2025-08-20	315d	B-10	C-08	\$39,360.00	DENIED_WITH_REBATE_PAID
2025-09-13	291d	A-11	C-00	-\$59,102.93	reopened from PENDING REVERSAL_CASH_NOT_RETURNED
2025-09-29	275d	A-13	C-00	\$1,197.16	reopened from PENDING RECOUPMENT_UNTRACEABLE INSUFFICIENT_DATA
2025-10-04	270d	A-06	C-05	\$4,149.00	reopened from PENDING SETTLEMENT_MISSING REBATE_AWAITING_MANUFACTURER

Queue table

Today. Heading EXCEPTION queue — the raw enum interpolated into a string, then uppercased by CSS. A Rank by dropdown, a count, then an 8-column table capped at 46vh with sticky headers.

Columns: Episode · Track · Service date · Age · Reimb. · Rebate · Variance · Why

A live row: E-000013 · Medical · 2025-08-20 · 315d · B-10 · C-08 · \$39,360.00 · reopened from PENDING DENIED DENIED_WITH_REBATE_PAID

Why it fails.

- EXCEPTION queue shouts a database word.
- Reimb. is an abbreviation with a full stop.
- The Reimb. and Rebate columns are two bare codes with no words anywhere on screen to decode them. B-10 means the insurer denied the claim and nobody has appealed. C-08 means the rebate was approved, paid, and the cash landed. Together they are the compliance story. On screen they are four characters each.
- The Why column is a bag of SCREAMING_SNAKE chips — up to five per row — with no legend. QUALIFICATION_UNDERMINED_BY_RECOUPMENT is a real live value.
- Why is the only column allowed to wrap while every other cell is nowrap, so row heights are ragged.

- Negative variance means *we owe a refund* — explained only in a hover tooltip.
- `315d` for age. Four different monospace treatments in one row.

Build this instead.

- Heading: `Claims that need work`, count as subtitle: `40 claims · $517,240.34 outstanding`.
- Merge `Reimb.` and `Rebate` into one column, `What's wrong`. Plain-English verdict name as primary text, code as muted monospace suffix, two stacked lines when both tracks have something to say:

Insurer denied the claim `B-10`
 Rebate paid in full `C-08`

- Cap `why` at two chips plus `+3 more`, each carrying its plain-English label from Part 4. The code goes in the hover.
- A cross-track flag must not look like an ordinary chip. `X-1` is a compliance exposure. Give it the `--status-serious` treatment CSS already defines, plus an explicit label: `Compliance – denied but rebate paid`.
- Reopened is a row state, not a chip. Move it out of `why` and mark the row itself — a left edge or a narrow dedicated column.
- Age: `315 days` or `10 months`. Not `315d`.
- Variance: keep the U+2212 minus and tabular figures. Put a visible legend above the table — `Negative = we owe a refund` — instead of a tooltip.
- Rename the `Rank by` options: `Reopened first, then oldest` → `Reopened first`; `Oldest first`; `Largest variance first` → `Biggest money first`; `Episode id` → `Claim number`.
- Give `Track` an icon or colour rather than repeating the words `Pharmacy / Medical` forty times.

Hard limits. Do not add sorts — the API allows exactly five (`reopened_first`, `age_desc`, `age_asc`, `variance_desc`, `episode_id`) and returns 422 for anything else. **Do not add filters** — there is no filter parameter beyond disposition and cursor. No payer filter, no drug filter, no date range, no search. Designing one is designing a backend change.

Available fields per row. `episode_id`, `track`, `date_of_service`, `ndc11`, `age_days`, `disposition`, `reimbursement_verdict`, `rebate_verdict`, `reimbursement_variance_cents`, `rebate_variance_cents`, `total_variance_cents` (signed), `absolute_variance_cents` (unsigned), `reopened_from`, `reason_codes[]`. `ndc11` and `absolute_variance_cents` are returned and **not displayed today**. There is no total-row-count field; `limit` caps at 1000 with no pagination.

1.5 Feed-level exceptions

3

OWS

OWS

Refresh

only by schema trigger, so there is no "done" state and nothing on this
ow here was committed on its own episode's analysis screen, behind the
ds what already happened. It re-reads itself when this tab regains focus,

le with "Inspect using AI" and commit a work item on its analysis screen.

372 reachable verdict pairs compositionally — the engine resolves each
otates, rather than enumerating combinations.

REBATE	EPISODES
C-00	2
C-00	2
C-08	1
C-00	1
C-05	1
C-07	1
C-01	1
C-09	1
C-10	1
C-11	1
C-14	1
C-01	1
C-14	1
C-00	1
C-05	1
C-13	1
C-00	1
C-01	1
C-05	1
C-08	1
C-13	1
C-01	1
C-02	1
C-08	1
C-00	1

Feed-level exceptions, default state

	WHY	RECEIVED
T	NO_KEYS_PRESENT	2025-08-28
S	NO_KEY_MATCH	2025-11-21
	NO_KEYS_PRESENT	2025-12-11
T	NO_KEYS_PRESENT	2025-12-17
	NO_KEY_MATCH	2025-12-24
	NO_KEYS_PRESENT	2025-12-29
T	NO_KEYS_PRESENT	2026-01-08
S	NO_KEY_MATCH	2026-01-14
T	NO_KEYS_PRESENT	2026-01-15
	NO_KEYS_PRESENT	2026-01-26
	NO_KEY_MATCH	2026-07-01
	NO_KEY_MATCH	2026-07-01
	NO_KEY_MATCH	2026-07-01
	NO_KEY_MATCH	2026-07-01
E	NO_KEY_MATCH	2026-07-01

WS

ent that cannot be tied to any qualified dispense in any other feed. Not a windfall — ce is audit exposure.

WS

l twice — an SFTP rerun, or a retry after an ambiguous timeout. Collapsed on the med: summing would double-count real money.

Refresh

By schema trigger, so there is no "done" state and nothing on this v here was committed on its own episode's analysis screen, behind the s what already happened. It re-reads itself when this tab regains focus,

with "inspect using AI" and commit a work item on its analysis screen.

72 reachable verdict pairs compositionally — the engine resolves each lates, rather than enumerating combinations.

REBATE	EPISODES
C-00	2
C-00	2
C-08	1
C-00	1
C-05	1
C-07	1
C-01	1
C-09	1
C-10	1
C-11	1
C-14	1
C-01	1
C-14	1
C-00	1
C-05	1
C-13	1
C-00	1
C-01	1
C-05	1
C-08	1
C-13	1
C-01	1
C-02	1
C-08	1
C-00	1

2025-09-04

VERDICT

TRANSITION
FIRST
REIMBURSEMENT VERDICT
B-02
REBATE VERDICT
C-00
EPISODE DISPOSITION
PENDING
REASON CODES
AWAITING_REMITTANCE

2025-09-12

Remittance claim line

CLP07
20260812689297
AMOUNT
\$0.00
CLP01
ENC-519314-00015
CLP02 CLAIM STATUS CODE
4
CHARGE
\$60,000.00

Feed-level exceptions, everything expanded

Today. Heading `Feed-level exceptions`, a two-line paragraph about the episode state space, then five independent `<details>` accordions. Each summary: a code chip, a bold title, a row count. Each body: a ~30-word paragraph and a three-column table.

ORDER	CHIP	TITLE	LIVE ROWS	DEFAULT
1	D-3	Orphan deposits	3	open
2	D-7	Allocation residuals	10	open
3	D-6	Parked records	15	collapsed
4	D-4	Unmatched rebates	0	collapsed
5	D-1	Duplicate deliveries	0	collapsed

Why it fails.


- `D-3`, `D-7`, `D-6` are unexplained and out of numeric order on screen. Nothing anywhere says what D means. It is a fourth code namespace — *feed-level data exceptions*, a separate axis from the A/B/C verdict tracks — and the screen never says so. A reader who has learned A, B and C are tracks will reasonably assume there is a track D. **There is not.**
- The open/closed rule is backwards. It opens a section when `0 < rows < 12`. So the section with the *most* problems (`D-6`, 15 rows) starts hidden, while two empty sections are also hidden — the panel can look busy while concealing its biggest item.
- Two ~30-word rationale paragraphs are visible simultaneously by default.
- Column labels unreadable alone: `Caused by`, `Bank row`, `Payload digest`, `ACH trace`.
- `Payload digest` prints 48 characters of a SHA-256 hash, silently truncated with no ellipsis.
- `record_kind`, `park_reason` and `source_system` all render in monospace as though they were identifiers.

Build this instead.

- Rename the panel to what it is: `Money and documents we could not match to a claim`.
- Lead with the dollars. `D-3` and `D-7` are dollar-denominated and cursor-aware. This panel holds real unattributed cash and currently presents it as a collapsed accordion.
- Sentence titles, code demoted:
 - `D-3` → `Cash we received but cannot attribute` · 3 items
 - `D-7` → `Deposits that did not fully allocate` · 10 items
 - `D-6` → `Documents whose ID did not match anything` · 15 items
 - `D-4` → `Rebate money we cannot tie to a dispense` · 0 items
 - `D-1` → `The same file delivered twice` · 0 items
- Reorder by consequence, not code number. Dollar-carrying first, then document-matching, then hygiene.
- Fix the disclosure rule. Open what has rows, collapse what is empty, never hide the biggest. Put an aggregate line in the panel header so nothing important sits behind a click.
- Keep the blurbs — they are excellent — but show one line each, or move them behind a `?`.
- Rename columns: `Caused by` → `From the deposit received on`; `Bank row` → `Bank line`; `Payload digest` → `File fingerprint`, truncated visibly with `...`.
- Preserve the distinction between the five. The narrative line *"four different problems with four different owners"* is exactly this panel. Do not merge them into one "Unmatched" list.

Available fields. `GET /api/feed-exceptions` → `orphan_deposits[]`{`norm_id`, `amount_cents`, `received_at`, `ach_trace_number`} · `allocation_residuals[]`{`allocated_cents`, `bank_norm_id`, `caused_by_received_at`} · `parked[]`{`norm_id`, `record_kind`, `park_reason`, `received_at`} · `orphan_rebates[]`{`norm_id`, `record_kind`, `received_at`, `park_reason`} · `duplicate_deliveries[]`{`source_system`, `payload_sha256`, `deliveries`}. UI caps rows at 40 with a `showing 40 of N` note. `duplicate_deliveries` ignores the cursor.

1.6 Portfolio Analyst



Portfolio Analyst

Today. A panel with `Portfolio Analyst`, an `Analyse the book` button, a five-line paragraph, and an optional question textarea. On run it streams a progress log, then renders four labelled prose blocks with inline citation superscripts.


Why it fails.

- It sits **second on the page, above the money**. An AI panel outranking the book's own numbers reads as a demo feature, not a product.
- The heading names an internal agent role. Nobody outside the team knows what a "Portfolio Analyst" is.
- The five-line paragraph is the densest copy on the screen and it explains architecture.
- Section labels come from raw response keys with underscores swapped for spaces: `state_of_the_book`, `what_is_concentrated`, `what_i_could_not_determine`, `where_to_look_first`.
- The 503 path shows the server's raw message. The long-run warning mentions "the provider's fast tier is down" — an internal operational detail.

Build this instead.

- **Move it below the worklist.** It answers a question asked *after* seeing the numbers.
- Heading: `Ask about the book`. Drop the role name.
- Reduce the paragraph to one line: *Every figure it quotes was computed by the engine and cited — it explains the numbers, it never produces one.* That sentence is the differentiator and it is currently buried at word 40 of 90.
- Real headings for the four output sections: `The state of the book` · `Where the money is concentrated` · `What I could not determine` · `Where to look first`.
- **Keep the citation superscripts and the unsourced-figure warning.** They are the proof in the narrative. Style the warning as a proper alert, not a bare outlined box.
- Style `What I could not determine` as a **feature**, not an apology — it is "the system never guesses" made visible.

1.7 To-do list



To-do list panel

Today. Heading `To-do list – all episodes`, a `Refresh` button, and a four-line hint reading, verbatim: *"Read-only: work_item is append-only by schema trigger, so there is no "done" state and nothing on this screen mutates anything..."* Empty state: *"No to-dos yet. Open an episode with "Inspect using AI" and commit a work item on its analysis screen."* Seven columns when populated: `Work item` · `Episode` · `Action` · `Rationale` · `Created` · `By` · `Artifacts`.

Why it fails.


- The hint leaks a **database table name** (`work_item`) and the phrase `append-only by schema trigger` to a business reader. It also uses a raw `--` instead of an em dash.
- The hint is longer than the table it describes.
- `recommended_action` renders in the same `.verdict-code` pill style as verdict codes, so two entirely different kinds of thing look identical.
- `created_by` prints a raw system identifier.

Build this instead.

- Heading: `Actions committed by the team`.
- One-line hint: *Every row here was approved by a person on a claim's own page. This list only reads — nothing here can be edited or checked off.*
- Give the action its own visual treatment, distinct from verdict codes — it is a decision, not a diagnosis. Label the seven values in plain English (Part 4).
- Format `created_by` as a person and keep it. *"The record of who did it names the person, not the model"* is a narrative point worth showing.
- Keep the "hidden by cursor" notice. Counting what the cursor hides rather than silently dropping it is the same honesty discipline as the rest of the page.

Available fields. `GET /api/agent/work-items?limit=` → `work_items[]`{`work_item_id`, `episode_id`, `recommended_action`, `summary`, `created_at`, `created_by`, `required_artifacts[]`, `at_cursor`}.

1.8 The episode side panel (right column)



Episode selected

Shared with Screen 2 and specified in full there (§2.2–2.5). On Screen 1 it is sticky, `3fr/2fr` , `max-height: calc(100vh - 48px)` , collapsing below the left column at 1100px.

Screen-1-specific issues:

- Its empty state is a 40-word paragraph. Replace with a short line plus a pointer.
- The `onClose` prop is destructured and never used — **there is no way to dismiss the panel** once an episode is selected.
- The `Inspect using AI` button only appears after a dossier loads, so the panel's most important affordance materialises late.

Screen 2 — The episode page (/analyse/{id})

E-000013

The agent type must complete a claimant every figure below has already been verified by the deterministic engine. The agent type explores E and recommends what a human should do next. — Back to the dashboard

EPISODE DOSSIER

The evidence the recommendation on the right has to answer is — the same object the agent type receives in a single call, so nothing about it can be checked anywhere else.

E-000013

BEVACJUNAR 400 MG/ML M.L. 400 units — medical benefit: BLUE HARBOR HEALTH

Accumulated: 2015-08-28 — LUMS: MC: 500000000

STATUS: [SELECTED BY AGENT](#) [EXCEPTION](#) [CASH](#) [CASH](#) — a story to single-track view can be set

[Status](#) [Details](#)

REIMBURSEMENT: \$6,000 of \$30,000 expected — 2015 REARS: \$15,000,000 of \$15,000,000 expected

WHAT HAPPENED: 2 RECORDS — 2 CASH MOVEMENTS — 5 VERDICT CHANGES ACROSS 11 EVALUATIONS

2015-08-28	Medical submission	STATUS: SELECTED BY AGENT												
2015-08-28	Service lines	<table> <tr> <th>CHARGE</th><th>ENCODING</th><th>COMPOSITE</th><th>UNITS</th></tr> <tr> <td>990,000,00</td><td>HC</td><td>99001-JW</td><td>4</td></tr> <tr> <td>990,000,00</td><td>HC</td><td>99001</td><td>1</td></tr> </table>	CHARGE	ENCODING	COMPOSITE	UNITS	990,000,00	HC	99001-JW	4	990,000,00	HC	99001	1
CHARGE	ENCODING	COMPOSITE	UNITS											
990,000,00	HC	99001-JW	4											
990,000,00	HC	99001	1											
2015-08-28	Medical acknowledgment	2015-08-28 ACCEPTED FOR PROCESSING												
2015-08-28	Cash acknowledgment	2015-08-28 STATUS: RECEIVED												
2015-08-28	Verdict	TRANSACTION: REIMBURSEMENT VERDICT RECORDED: B-00 EPISODE DISPOSITION: PENDING PREVIOUSLY CLOSED AT: 2015-08-07/23/28-59C												
2015-08-28	Reimbursement claim line	CLAIM CLINE STATUS: CLOSED ADJUSTMENTS: \$60,000,000 AMOUNT: \$60,000,000 GROUP CODE: RARC REASON CODE: R0												
2015-08-28	Verdict	TRANSACTION: REIMBURSEMENT VERDICT RECORDED: B-00 EPISODE DISPOSITION: PENDING PREVIOUSLY CLOSED AT: 2015-08-07/23/28-59C												
2015-08-28	Beacon payment reference	BEACON ACCOUNT: 20000000 FINANCIAL EFFECTIVE DATE: 2015-0000												
2015-08-28	Beacon validation outcome	VALIDATION OUTCOME: ACCEPTED												
2015-08-28	Cash	DIRECTION: TRADE IN: ALLOCATION CODE												
2015-08-28	Rebate dispute line	STATUS: SELECTED BY AGENT REBATE AMOUNT: \$15,000,000												
2015-08-28	Verdict	TRANSACTION: REIMBURSEMENT VERDICT RECORDED: B-00 EPISODE DISPOSITION: PENDING PREVIOUSLY CLOSED AT: 2015-08-07/23/28-59C												
2015-08-28	Type qualification	STATUS: SELECTED BY AGENT QUALIFIED												
2015-08-28	Verdict	TRANSACTION: REIMBURSEMENT VERDICT RECORDED: B-00 EPISODE DISPOSITION: PENDING PREVIOUSLY CLOSED AT: 2015-08-07/23/28-59C												

1. EXPLAN


The Decision Investigator reads this episode's dossier and writes a narrative (using coding records, engine events and verdicts by it), never inventing a figure of its own.

2. DECIDE NEXT STEPS

The Decision Investigator presents an action, an independent Evaluator grades it a 0-5-point verdict, and the top repeats until the score clears the line, the proper one to stop, or a budget runs out. Every step takes in the last record's own audit trail, otherwise line — nothing here is a summary written after the fact.

Episode page, full


Today. Two columns. Left: the same episode dossier component as Screen 1. Right: two stacked panels, `1. Explain` and `2. Decide next steps`. Masthead is a bare `<h1>` reading `Agent analysis – E-000013`, a paragraph, and a plain text link `← Back to the dashboard`.



Episode page, above the fold

Structural note. This page is reached only by a full browser navigation into a new tab. There is no router and no in-app back. It is, in effect, a second application sharing a stylesheet.

2.1 Episode page masthead



Episode page masthead


Today. `<h1>Agent analysis — E-000013`. Subtitle: *"The agent layer never computes a number: every figure below was already decided by the deterministic engine. The agent only explains it and recommends what a human should do next."* Then an inline anchor `<a href="#">← Back to the dashboard`.

Why it fails. The page is titled after the *technology* rather than the *claim*. An executive opening this wants to see the claim.

Build this instead.

- `<h1>` carries the claim's human identity — `BEVACIZUMAB 400 MG • Blue Harbor Health • dispensed 20 Aug 2025` — with `E-000013` as a muted monospace suffix.
- Put status and outstanding dollars in the masthead so they survive scrolling: `$39,360.00 outstanding • needs work`.
- Make `<a href="#">← Back` a real button, top-left, not an inline link inside a paragraph.
- Keep the "never computes a number" sentence. Shorten it, set it as a quiet line, not a lead paragraph.

2.2 Dossier header



Dossier header close-up

Today, in render order: episode id at 16px mono/650 · a composed drug line (BEVACIZUMAB 400 MG/16 ML · 400 units · medical benefit · BLUE HARBOR HEALTH) · a keys line (dispensed 2025-08-20 · CLM01 ENC-519314-00015) · right-aligned: EXCEPTION in raw uppercase, then two unlabelled code pills B-10 C-08 , then reopened from PENDING · then chips DENIED DENIED_WITH_REBATE_PAID X-1 followed by loose text "cross-track — a story no single-track view can tell" · then a Simple / Detailed toggle.


That is 17+ distinct facts in four font sizes and three families before the timeline starts.

Why it fails.

- The two verdict pills are the worst single element in the product. Two opaque codes side by side with nothing saying which is reimbursement and which is rebate. The only place that mapping is stated is a column header inside a collapsed accordion further down the page.
- EXCEPTION is a raw enum. CLM01 is an EDI segment identifier printed as a field label.
- The X-1 chip looks identical to the two beside it, but it is a compliance exposure and they are ordinary reason codes.
- "cross-track — a story no single-track view can tell" is a pitch line sitting in a data region, rendering as loose text after three chips.

Build this instead.

- Two labelled rows, not two bare pills:




- Status as a word with its colour: Needs work , EXCEPTION as a muted suffix.
- reopened from PENDING → Was settled on 1 Sep 2025, reopened since . The date is already in the payload as previously_closed_at .

- Give the cross-track flag its own band — `--status-serious`, full width, with a sentence: "Compliance exposure — the insurer refused to pay for this drug and the manufacturer paid the 340B rebate anyway. Net position is negative." This is the best story in the product and it currently renders as a 3-character chip.
- Hide `CLM01`, `Rx`, `fill_number` and NPIs behind a `Claim identifiers` disclosure. An executive never needs them; an operator needs them on demand.
- Keep the `Simple / Detailed` toggle. It works, and the split is computed server-side so the AI layer and the UI agree.


Available fields. `GET /api/episode/{id}/dossier` → `identity{track, drug, ndc11, quantity, date_of_service, payer, pharmacy_npi, rx_number, fill_number, clm01, billing_provider_npi, is_340b_flagged, opened_at}` · `current{reimbursement_verdict, rebate_verdict, episode_disposition, reason_codes[], cross_track_flags[], reopened_from, as_of}` · `previously_closed_at` and the per-track `reimbursement_disposition` / `rebate_disposition` exist on `GET /api/episode/{id}` and are shown nowhere today.

2.3 The money band



Money band

Today. A grey rounded band with two cells:



Why it fails. This is the most important region on the page for the target audience, styled as a subdued utility strip in 13px text. `$39,360.00 outstanding` — the entire reason the claim is in the queue — is the same size and weight as the words around it.


Build this instead.

- Make this the visual anchor of the page, directly under the header. The outstanding figure gets display treatment: 28px, tabular, `--status-critical` when money is owed to us.


- Three explicit labels per track, not a run-on sentence: `Expected` / `Received` / `Outstanding` .
- Colour the rebate track with `--track-rebate` (`#4a3aa7`) — the token exists and the stylesheet never uses it.
- When the rebate is absent the cell prints `no rebate expected` . Suppress the cell rather than showing an apology where a number should be.
- Show the net position when the two tracks disagree in sign. E-000013 is exactly that case: \$39,360 owed to us on one track, \$15,000 already received on the other, on a claim the insurer denied. That net is the compliance story.

Available fields. `economics{expected_reimbursement_cents, received_reimbursement_cents, reimbursement_variance_cents, expected_rebate_cents, received_rebate_cents, rebate_variance_cents}` . All integer cents. **No percentage field exists anywhere in the API** — any ratio must be computed in the browser.

2.4 The timeline



Timeline, Simple view



Timeline, Detailed view

Today. An ordered list in a `130px 1fr` grid. Left gutter: right-aligned monospace `YYYY-MM-DD` and a coloured status dot. Right: a 14px bold headline, an optional `late-arriving` pill, then a `repeat(auto-fill, minmax(180px, 1fr))` grid of label/value pairs. Detailed view adds a provenance line with a clickable `file:line` link expanding raw JSON and two SHA-256 hashes. Heading above: `What happened` followed by `8 records · 2 cash movements · 4 verdict changes across 11 evaluations`.

Why it fails.


- `humanizeTag` destroys acronyms. It lowercases the whole tag then capitalises only the first letter, producing `Tpa qualification`, `Tpa rebate request`, `Tpa manufacturer decision`, `Tpa reversal` on screen.
- `humanizeKey` uppercases anything matching a short letter+digit pattern, so raw EDI segment ids render as field labels: `CLP02 claim status code`, `SVC01 composite`, `STC12 free form`.
- Fact keys are never curated — whatever the payload holds becomes a `<dt>`. A `VERDICT` event shows five rows of enums: `Transition: FIRST`, `Reimbursement verdict: B-02`, `Rebate verdict: C-00`, `Episode disposition: PENDING`, `Reason codes: AWAITING_REMITTANCE`.
- Tables render **inside a definition list** for nested line items.
- The counts string welds four numbers onto a heading.
- **The status dot is clipped** — Part 5, bug 1.

Build this instead.

- **Fix the two humanisers.** Use an explicit label map for the 16 event tags rather than string-mangling:
 - `TPA_QUALIFICATION` → **340B administrator ruled on eligibility**
 - `TPA_MANUFACTURER_DECISION` → **Manufacturer ruled on the rebate**
 - `REMITTANCE_CLAIM_LINE` → **Payer's payment decision**
 - `PROVIDER_LEVEL_ADJUSTMENT` → **Payer clawback**
 - `BANK_TRANSACTION` → **Money moved at the bank**
 - `MEDICAL_ACKNOWLEDGMENT` → **Clearinghouse accepted the claim**
 - `REBATE_DISPENSE_LINE` → **Rebate line on a manufacturer payment**
- **Group the timeline into three lanes or three colours** — insurance, 340B rebate, cash. That parallelism is the whole domain insight and a single column hides it.
- **VERDICT events should read as a sentence, not a field list:** *"1 Oct 2025 — reopened. Now: denied by the payer. Was: awaiting payment."*
- Long dates: `20 Aug 2025`.
- Keep `late-arriving` and the `happened X, learned Y` treatment. Late arrival is a core domain fact, already well handled.
- Keep the provenance link and the raw-record panel. Being one click from the literal source line is the proof in the narrative — it deserves to be *more* visible in Detailed view, not less.
- Move the counts line off the heading into its own quiet row.

Available fields. `timeline[] {at, occurred_on, tag, facts{}, essential[], source{}, late} · counts {records, cash_allocations, verdicts_written, verdict_changes}`. `essential[]` is a server-authored list of which fact keys matter — the Simple/Detailed split is already in the payload.

2.5 Evidence accordion and unresolved records



Evidence accordion expanded

Today. A closed-by-default `<details>` summarised "The raw material — every verdict written, and every key that resolved here", containing two stacked tables:

- **Verdict log** — `Evaluated at` · `Disposition` · `Reimb.` · `Rebate` · `Reasons`. The disposition cell is a literal `●` glyph followed by the raw enum.
- **Crosswalk keys** — `Key type` · `Value` · `First seen`.

Above it, when present, an **unresolved records** table — `Record` · `Why it parked` · `Key it carried` · `Arrived` — under "Records that tried to reach this claim and could not".

Why it fails.


- `key_type` chips are raw: `NCPDP_CLAIM`, `MEDICAL_CLM01`, `PAYER_ICN`, `TRN02`, `ALLOCATION_CODE`, `NATURAL_340B_PHARMACY`, `NATURAL_340B_MEDICAL`, `PBM_AUTH`, `BEACON_ID`, `COVERED_ENTITY_340B`, `HCPCS`, `PAYMENT_REFERENCE`.
- `park_reason` chips are raw: `NO_KEY_MATCH`, `AMBIGUOUS_KEY_MATCH`, `NO_KEYS_PRESENT`, `COVERED_ENTITY_MISMATCH`.
- The `●` glyph is a second, inconsistent treatment of the same status-dot concept used in the timeline.
- The verdict log is where a reader would finally learn that the first code is reimbursement and the second is rebate — and it is collapsed by default.

Build this instead.

- Split into two disclosures with plainer summaries: `How this claim's answer changed over time` and `Which identifiers matched this claim`.
- The verdict log is the best audit artefact in the product. Surface a one-line summary outside the accordion: "Evaluated 11 times, the answer changed 4 times." Both numbers are already in `counts`.

- Give `park_reason` its plain-English label (Part 4). These are the four "different problems with different owners" from the narrative.
- Keep the unresolved-records table prominent. Records that tried to reach this claim and failed are by definition invisible everywhere else. Heading: `Documents that may belong to this claim but did not match`.

2.6 Explain and Decide



Explain and Decide, idle

Today, idle. Two panels: 1. Explain with an Explain button and a two-line description; 2. Decide next steps with a disabled button (tooltip Run Explain first) and a four-line description.

On run, Explain produces four prose blocks labelled from raw keys — what happened , why it is open , what i could not determine , what a human should check first — plus a citation count and an optional unsourced-figure warning.

On run, Decide streams an append-only log into a 42vh scroll box, one row per event, dividers reading Iteration 1 , Iteration 2 . Event kinds render as tool_call {name}({args truncated to 140 chars}) , ↳ {status} ({error_type}) · {bytes} B · {elapsed_ms} ms , score: 0.847 — 2 criteria not supported , gate: continue (score 0.847) . Then an outcome banner and an editable work-item form.

What to keep. The 1. / 2. numbering and the disabled second button are good product design. The gating rule — nobody should read a recommendation before reading what happened — is a real safety property. Make it explicit in the UI rather than hiding it in a tooltip.

Why it still fails.

- The idle state is almost entirely empty space with two paragraphs of architecture.
- The live log leaks internal event names to screen: llm_error: , tool_call , forced_tool_name_coerced , injection_attempt_recorded , plus an unhandled-kind fallback printing {kind}: {raw JSON} .
- Tool results are never shown — only status, error type, byte count and elapsed ms. The evidence the report cites never reaches the screen.

- `proposal.evidence[]` and `result.citations[]` are fetched and **only their `.length` is rendered**.
- Sixteen graded criteria with written rationales exist per run; the UI shows a chip and a `title` tooltip.
- `criterion_id` is truncated to the pre-underscore token in the stream but shown in full in the banner — same data, two renderings.
- A known confusing state: three green `SUPPORTED` chips can sit directly above `score: 0.000` with nothing explaining it, because the four vetoes are deterministic and never appear in the evaluation event. **It reads as a broken scorer rather than as a veto working.**

Build this instead.

- **Two named steps with a visible dependency:** `Step 1 – Understand what happened` → `Step 2 – Decide what to do`. Show the lock reason on the face of step 2.
 - Real headings for the four Explain sections, properly capitalised.
 - **The log is the product, so make it legible rather than raw.** Group by iteration into collapsible blocks. Translate the event kinds: `tool_call` → `Looked up: the claim's payment records`; `lock · 4.2 KB · 310 ms` → `Found 8 records`. Keep a `Show raw log` toggle for the engineer in the room.
 - **When a veto fires, say so.** A `score: 0.000` beneath green chips needs one line: *"Blocked by a hard check: a figure appeared that no tool returned."* Style a veto distinctly from a low score.
 - Surface `proposal.evidence[]` and the citation list. Already on the wire, and they are the reason to trust the output.
 - Move criterion rationales out of `title` attributes into an expandable list.
 - **Keep the outcome vocabulary and style `Insufficient data` as a success.** Four outcomes exist: `complete`, `insufficient_data`, `stalled`, `capped`. The claim *"I could not tell" is a first-class answer* depends on this being visible.
-

2.7 The human gate (work-item form)


Today. When Decide produces a proposal, a form appears: `Proposed work item – edit before committing`, an `Action` select with seven options, a `Rationale` textarea pre-filled from the model's reasoning with a live `{n}/600, min 20` counter, and one text input per required artifact with add/remove buttons. Then `Add to-do`, which runs a two-call commit.

Why this matters more than any other element on the page. Every field is editable by design. The reasoning is recorded explicitly: a reviewer who can only approve or reject is *ratifying*; one who can correct the action is *judging*. The design is a direct response to documented industry failures where a human-in-the-loop step became rubber-stamping.

Build this instead.

- **Make the editability visible, not incidental.** Heading: `Your decision`, with a line: *"This is a draft. Change anything — the record names you, not the model."*
 - Show the model's original proposal alongside the editable fields, visibly marked as the draft, so a change reads as a change. The commit already sends a `proposed` object recording what the model said.
 - Label the seven actions in plain English (Part 4).
 - The success box reads `Work item created: {id}. {message}`. Make it a sentence.
 - Keep the 20–600 character rationale rule and the counter. A forced rationale is an anti-rubber-stamp mechanism.
-

Screen 3 — Verdict library (new screen)



Verdict distribution today

Today. The last panel on the landing page. Heading `Verdict distribution`, a two-line paragraph about 372 reachable pairs and 50 deterministic rules, then a three-column table: `Reimbursement` · `Rebate` · `Episodes`. Twenty-five rows of bare code pairs.

Live content: `A-04 / C-00 → 2`, `A-12 / C-00 → 2`, then **twenty-three consecutive rows with a count of 1**.

Why it fails, bluntly. It is the only table on the landing page with **no height cap** — it grows unbounded and can push the page to thousands of pixels. It carries **no money at all**, only episode counts. It is capped at 25 rows by the API (`LIMIT 25`), so its counts do not sum to the episode total and nothing on screen says so. And on the live demo profile, 23 of its 25 rows say "1". An executive reading it learns nothing, twice.

Cost of this change — read before building

This is **the only change in this document that is not purely cosmetic**. The nav is a two-value `useState`, not a router. A third screen means: one more value in that state, one more button in the segmented control, one more branch in the body render. No new endpoint, no new route, no new fetch — the screen consumes `overview.verdict_pairs[]`, already fetched on every page load. **Scope: one state value, one button, one conditional. Do not let this grow into routing work.**

Build this instead

- Name it for the question it answers: `What kinds of problems are in the book?`
- Stop showing raw pairs as the primary object. The pair table is a reference artefact; the executive view is the *families*. Group by what the operator does about it:
- Underpaid or short-paid
- Denied or rejected

- Paid but no cash arrived
- Rebate approved and unpaid
- Reversed, clawed back, or duplicated
- Compliance exposure (cross-track)
- Settled, nothing to do
- **Each family gets an episode count and a plain-English sentence.** Counts come from `overview.reason_codes[]` — already returned, **no row limit**, ordered count-descending. A strictly better source than `verdict_pairs` for this purpose.
- **Keep the pair table below as a reference section** — both codes carrying their plain-English label, plus an explicit note: *"Top 25 pairs by frequency. Counts do not sum to the book total."*
- **State the D-code namespace on this screen.** This is the natural home for the sentence that resolves the confusion: *"A and B verdicts describe the insurance track. C verdicts describe the 340B rebate track. X flags mark a problem visible only when you look at both. D exceptions are a separate thing entirely — faults in the data we received, not facts about any claim."*
- Include the Part 4 dictionary here. It is a reference screen; a legend belongs on it.
- **Cap the table height** with the same `max-height` + `overflow-y` treatment the queue table already uses.


Available fields. `overview.verdict_pairs[]{reimbursement, rebate, episodes}` (LIMIT 25, no money) ·

`overview.reason_codes[]{reason_code, episodes}` (no limit, count-descending) · `overview.cross_track_flags[]{flag_code, episodes}` (alphabetical, not by size).

None of these carry dollars. Money is aggregated by disposition only. A "top categories by dollars" view would have to be built client-side from queue rows — `absolute_variance_cents` plus `reason_codes[]` per row, three calls, a 1000-row cap, and no total count to tell you whether you got everything. That is an engineering decision, not a design one, and it is **out of scope under the freeze**.

Screen 4 — Connectivity

[illegible]



Connectivity, above the fold

Today, top to bottom.

1. An honesty banner: Live vendor connections: none. plus a generated statement and notes.
2. Where each source actually is — a four-rung ladder with a count per rung: Declared · Connector-ready · Working connection · Production-ready. Chips read the ceiling for this build and no source reaches this.
3. A 12-column gate evidence table: Source · Vendor · Stage · Reaches · Transport · Auth · Schema · Mock fidelity · Golden · Control totals · Gaps · Blocked.
4. A repository-level evidence panel and an archetype table.
5. A right-hand detail panel showing one selected source in full.

What is good here and must survive. This is the most product-mature screen in the build. It states what is *not* connected as loudly as what is. The banner sentence is generated from code at request time. The ladder refuses to say "in progress", with a stated reason: "a report that could say 'mostly' would always say 'mostly'." For an audience that evaluates delivery risk professionally, this is the most credible screen in the product. **Do not soften it.**

Why it still fails an executive.

- **Twelve columns.** Nine are engineering evidence: transport class, auth resolver, schema contract version, mock fidelity tiers, golden-claim traces, control totals.
- SPEC / STANDARD / INVENTED fidelity tiers are meaningful and completely unexplained on the face of the table.
- The rung labels are internal vocabulary. Connector-ready means *built and tested against a mock, switchable to live by config*. Working connection means *the vendor has issued us a credential*. Neither is self-evident.
- Row selection sits on the <tr> with no keyboard affordance and no visible button.
- generated_at prints as a raw unformatted timestamp.
- gaps[] and blocked[] collapse to integer counts; you must click a row to see any of them. There is **no aggregate** — no total gap count, no cross-source view of which credential would unblock the most sources, though all of it arrives in one response.

Build this instead.

- **Lead with the ladder, not the banner.** The four rungs with counts are the executive summary. Make them large.
- **Two columns by default** — `Source` and `Where it stands` (the rung plus one sentence) — expanding to the full evidence set behind a `Show evidence` toggle. Everything else moves to the detail panel, which already holds it all.
- Rename the rungs for a business reader, internal name as suffix: `Registered only` · `Built and tested, not yet authorised` · `Live connection` · `Production-hardened`.
- **Explain the fidelity bar inline.** It is already dual-encoded — colour plus text counts, so colour never carries meaning alone. Good. It needs one line: *"How much of this mock comes from the vendor's own documentation, versus an industry standard, versus our own invention."*
- **Add the cross-source rollup the data already supports** — one line above the table: *"14 sources registered · all connector-ready · 0 authorised · blocked on N vendor credentials."* Every number is in the one response.
- Format `generated_at` readably. Make row selection a real focusable control.

Available fields. `GET /api/connectivity` → `live_vendor_connections[]` · `live_vendor_connection_statement` · `notes[]` · `evidence_entries` · `unavailable_evidence[]` · `golden_test_module` · `generated_at` · `archetype_traces[]` · `sources[]` {`source_id`, `vendor`, `enabled`, `reaches`, `endpoint`, `stage`, `transport{}`, `auth{}`, `schema{}`, `mapping{}`, `fidelity{}`, `mock_modules[]`, `golden{}`, `control_totals{}`, `gaps[]`, `blocked[]`} · `active_tpa_source{source, datasets[], statement}` — which reports what was *read*, kept deliberately separate from what was *reached*.

Part 4 — The plain-English dictionary

This is the highest-value part of the redesign. The screens are currently readable only by someone holding a 43-entry lookup table. Nobody in the room will be.

Two rules, repeated because they matter:

- 1. **The word is the primary text. The code is a muted monospace suffix.** Never delete the code — an operator cites it to a payer, and the AI layer grounds on it.
- 2. **Never rename a code string in the data.** Codes feed the AI layer's grounding tables and recorded evaluation fixtures. A rename silently breaks the AI layer at the next rebuild.

Short labels already exist in code at `src/recon/domain/verdicts.py:126` (`VERDICT_DESCRIPTIONS`), reachable via `describe()` at line 142. The frontend has never imported them.

4.1 The four namespaces — say this on Screen 3

There are **four** code families on screen and only two of them are tracks. This is the confusion to kill:

FAMILY	WHAT IT DESCRIBES	COUNT
A and B	The insurance track . A = pharmacy benefit (PBM). B = medical benefit (insurer). These are two halves of ONE track , not two tracks — a claim travels one road or the other, never both.	16 + 15
C	The 340B rebate track . An optional second track on the same claim.	12
X	Cross-track flags . A problem visible only when you look at both tracks together. Annotations, never prohibitions.	7
D	Feed-level data exceptions . A fault in the data we received — not a fact about any claim. There is no track D .	7

Every claim always carries **exactly two verdicts**: one insurance, one rebate. `C-00` is the code for "this claim has no 340B rebate", so the rebate slot is never empty. That is why they are always shown as a pair.

4.2 Insurance track — pharmacy benefit (A codes)

CODE	PLAIN-ENGLISH LABEL	WHAT THE OPERATOR DOES
A-01	Refused at the counter	Nothing. The drug never went out.
A-02	Waiting for payment	Wait.
A-04	Paid in full and settled	Nothing.
A-05	Paperwork says paid, no money arrived	Chase the payer for cash it says it sent.
A-06	Money arrived, closing confirmation missing	Chase the confirmation, not the cash.
A-07	Underpaid, and the bank agrees with the short figure	Dispute the shortfall.
A-08	Underpaid and no money arrived	Two faults. Chase both.
A-09	Overpaid — we owe money back	Book a refund liability.
A-10	Cancelled, money returned	Nothing — but check no rebate still stands on it.
A-11	Cancelled, but we still hold the cash	Our books overstate revenue. Return it.
A-12	Clawed back after audit, and we traced it	Nothing.
A-13	Clawback we cannot tie to any bank movement	Investigate. The engine declines to decide.
A-14	Contract explains the whole gap	Nothing. Not a shortfall.
A-15	Contract explains part of the gap, not all	Chase the unexplained remainder.
A-16	Cancelled before any money moved	Nothing.
A-17	Paid twice for one claim	Refund one.

4.3 Insurance track — medical benefit (B codes)

CODE	PLAIN-ENGLISH LABEL	WHAT THE OPERATOR DOES
B-01	Bounced by the clearinghouse	The insurer never saw it. Fix and resubmit.
B-02	Submitted, waiting on the insurer	Wait.
B-04	Paid in full, cash matched	Nothing.
B-05	Insurer says paid, no money arrived	Chase the money.
B-06	Paid short, no appeal filed	Decide: appeal, or write off.
B-07	Paid short, appeal in progress	Wait.
B-08	Appeal won, balance paid	Nothing.
B-09	Appeal lost	Make a write-off decision.
B-10	Denied, no appeal filed	Decide: appeal, or write off.
B-11	Denied, appeal in progress	Wait.
B-12	Denial overturned, paid	Nothing.
B-13	Denied, appeal lost	Total loss. Check upstream eligibility screening.
B-14	Appeal won, payment never arrived	The highest-value chase in the system.
B-15	Paid, then the insurer took it back	Investigate and dispute.
B-16	The same remittance arrived twice	De-duplicate before it reaches the ledger.

4.4 340B rebate track (C codes)

CODE	PLAIN-ENGLISH LABEL	WHAT THE OPERATOR DOES
C-00	No 340B rebate on this claim	Nothing. Absence, not a problem.
C-01	Waiting on the eligibility ruling	Wait.
C-02	Ruled not eligible	Nothing. A correct answer, not a failure.
C-03	Eligible, rebate not yet requested	Wait, or push the submission.
C-05	Requested, manufacturer has not answered	Wait.
C-07	Manufacturer refused the rebate	Decide whether to correct and resubmit.
C-08	Approved, paid, cash matched	Nothing.
C-09	Administrator says paid, bank shows nothing	Chase the rebate cash.
C-10	Paid less rebate than expected	Dispute the shortfall — usually a unit or price dispute.
C-11	Approved, waiting on payment	Wait.
C-13	A rebate we received was reversed	Investigate — the original eligibility is now contested.
C-14	Two rebates paid for one dispense	Refund one.

The numbering has deliberate holes. A-03 , B-03 , B-17 , C-04 , C-06 , C-12 are retired and asserted absent in code. Do not render a continuous range — it would show codes that cannot exist.

4.5 Cross-track flags (X codes) — these are compliance, not diagnostics

Give these a distinct visual treatment. They are the reason a two-track view exists at all, and today they render identically to ordinary reason chips.

CODE	PLAIN-ENGLISH LABEL	FORCES EXCEPTION?
X-1	Insurer refused to pay, but we kept the 340B rebate. Net position negative; audit exposure.	Yes
X-2	The sale was cancelled but a rebate was paid anyway. The drug never reached the patient, so the rebate is invalid.	Yes
X-3	Money clawed back for ineligibility, and the rebate rests on the same now-wrong facts.	Yes
X-4	Insurance is clean; only the rebate was refused. Deliberately not a compliance flag.	No
X-5	Both tracks show "paid on paper, no money in the bank". Almost certainly one cause on our side, not two payers failing at once.	Yes
X-6	Every channel failed. Total loss on the drug.	Yes
X-7	An appeal was won after the rebate was clawed back — the reason for the clawback may no longer hold. Recoverable money nobody is watching.	Yes

X-4 exists and does nothing on purpose: the engine must recognise the shape in order to *not* escalate it.

4.6 Feed-level data exceptions (D codes) — the fourth namespace

There is no track D. These describe faults in the data we received, orthogonal to any claim's state.

CODE	CURRENT UI TITLE	PLAIN-ENGLISH LABEL
D-1	Duplicate deliveries	The same file delivered twice. An SFTP rerun or a retry. Counted once, never summed — summing would double-count real money.
D-2	<i>(no code shown; renders as an "Arrival lagged the event" chip)</i>	A record arrived after the thing it explains. Not an error — it becomes matchable later.
D-3	Orphan deposits	Cash we received but cannot attribute. Not a windfall — cash held without provenance is audit exposure.
D-4	Unmatched rebates	Rebate money we cannot tie to any dispense in any feed.
D-5	<i>(absent)</i>	Malformed record. Deliberately never generated. Correctly not on screen.
D-6	Parked records	A document whose ID matched nothing. Held, not discarded, and re-checked on every later arrival. <i>The claim is not missing — the mapping failed.</i>
D-7	Allocation residuals	A deposit that did not fully allocate. The deposit is right; our split of it is incomplete.

One naming inconsistency to resolve in the design. The UI files the "Parked records" panel under `D-6`, but the API docstring serving that endpoint names only D-1, D-3, D-4 and D-7. A *parked record* is the mechanism; `D-6` is the exception that mechanism surfaces. They are related, not synonyms. Pick one framing on screen and stay with it.

4.7 Why a document did not match (park reasons)

These are the narrative's "*four different problems with four different owners*". Keep them distinct.

CODE	PLAIN-ENGLISH LABEL
<code>NO_KEY_MATCH</code>	The identifiers it carried matched nothing we hold
<code>AMBIGUOUS_KEY_MATCH</code>	It matched more than one claim — we will not guess which
<code>NO_KEYS_PRESENT</code>	It carried no usable identifier at all
<code>COVERED_ENTITY_MISMATCH</code>	It matched cleanly but names a different 340B entity

4.8 Reason codes seen in the live demo data

The engine has 37. These are the ones on screen right now. **No authored label table exists for these anywhere in the repo** — their meaning is defined only by which verdict maps to them. The labels below are written for this brief and need one review pass by someone who knows the engine.

CODE	PLAIN-ENGLISH LABEL
AWAITING_REMITTANCE	Waiting for the insurer's payment decision
UNDERPAID	Paid less than the contract allows
OVERPAID	Paid more than the contract allows
NO_CASH	The paperwork says paid; no money arrived
CASH_MISMATCH	The money that arrived does not match the paperwork
SETTLEMENT_MISSING	Money arrived, closing confirmation did not
DENIED	The insurer refused to pay
REJECTED_AT_POS	Refused at the counter, before dispensing
CLEARINGHOUSE_REJECTED	Bounced before the insurer ever saw it
APPEAL_LOST	The appeal failed
APPEAL_WON_NO_CASH	The insurer agreed it owes us and still has not paid
ADJUSTMENT_RESIDUAL	The contractual discount explains part of the gap, not all
DUPLICATE_PAYMENT	Paid twice for one claim
DUPLICATE_REMITTANCE	The same remittance document arrived twice
REVERSAL_CASH_NOT_RETURNED	The sale was cancelled but we still hold the money
RECOUPMENT_UNTRACEABLE	A clawback we cannot tie to any bank movement
TOTAL_LOSS	Every collection path is exhausted
INSUFFICIENT_DATA	The engine deterministically cannot decide. Not an error — a first-class answer
REBATE_AWAITING_QUALIFICATION	Waiting on the 340B eligibility ruling
REBATE_AWAITING_MANUFACTURER	Requested; the manufacturer has not answered
REBATE_AWAITING_PAYMENT	Approved; the manufacturer has not paid
REBATE_NOT_QUALIFIED	Ruled not eligible for 340B
REBATE_NO_CASH	The administrator says paid; the bank shows nothing
REBATE_UNDERPAID	Less rebate paid than expected
REBATE_CLAWED_BACK	A rebate we received was reversed
DUPLICATE_REBATE	Two rebates paid for one dispense
REBATE_ON_UNDISPENSED_CLAIM	A rebate stands on a claim where the drug never went out
DENIED_WITH_REBATE_PAID	The insurer refused, the manufacturer paid anyway
QUALIFICATION_UNDERMINED_BY_RECOUPMENT	A clawback undercuts the facts the rebate was granted on
REBATE_RE_REQUEST_AVAILABLE	The rebate can be requested again

4.9 The seven actions a human can commit

CODE	PLAIN-ENGLISH LABEL
RESUBMIT	Fix and send the claim again
APPEAL	Challenge the insurer's decision
WRITE_OFF	Accept the loss and close it
ESCALATE	Hand to someone with authority to chase a counterparty
INVESTIGATE_CROSSWALK	The money may be here — the ID match failed
AWAIT_PAYER	Nothing to do yet; the other side owes us a response
ABSTAIN	Not enough evidence to recommend anything

4.10 Other raw strings currently on screen

WHERE	RAW STRING	SUGGESTED LABEL
Queue table	Reimb.	Insurance
Dossier header	EXCEPTION / PENDING / CLOSED	Needs work / Waiting / Settled
Dossier header	CLM01 ENC-519314-00015	Hide behind a Claim identifiers disclosure
Timeline	Tpa qualification	340B administrator ruled on eligibility
Timeline	CLP02 claim status code	Payer's status code
Timeline	STC12 free form	Clearinghouse note
Timeline	SVC01 composite	Procedure code
Feed exceptions	Caused by	From the deposit received on
Feed exceptions	Payload digest	File fingerprint
Feed exceptions	Bank row	Bank line
To-do list	work_item is append-only by schema trigger	Remove entirely
Crosswalk keys	NATURAL_340B_PHARMACY etc.	Pharmacy natural key etc.
Agent log	tool_call, llm_error, forced_tool_name_coerced	Translate; keep raw behind a toggle
Connectivity	SPEC / STANDARD / INVENTED	From vendor docs / From an industry standard / Our own invention
Connectivity	CONNECTOR_READY	Built and tested, not yet authorised

Appendix A — Screenshot index

All images are in `docs/product/img/`, captured live at 1440×900 on the `demo` profile, seed `8984883`, cursor `2026-07-01`.

FILE	WHAT IT SHOWS
<code>01-dashboard-full.png</code>	Operations, entire page
<code>02-dashboard-fold.png</code>	Operations, what loads above the fold
<code>03-masthead.png</code>	Masthead
<code>04-replay-cursor.png</code>	Replay cursor panel
<code>05-queue-tiles.png</code>	The three disposition tiles
<code>06-queue-table.png</code>	Exception queue table
<code>07-portfolio-analyst.png</code>	Portfolio Analyst panel, idle
<code>08-feed-exceptions.png</code>	Feed-level exceptions, default state
<code>09-verdict-distribution.png</code>	Verdict distribution table, all 25 rows
<code>10-todo-panel.png</code>	To-do list, empty state
<code>11-episode-selected.png</code>	Queue plus episode panel together
<code>12-dossier-header.png</code>	Dossier header close-up
<code>13-dossier-money.png</code>	Money band close-up
<code>14-timeline-simple.png</code>	Timeline, Simple view
<code>17-timeline-detailed.png</code>	Timeline, Detailed view
<code>18-dossier-full.png</code>	Whole dossier panel
<code>19-feed-exceptions-expanded.png</code>	Feed exceptions, all five expanded
<code>20-evidence-accordion.png</code>	Verdict log and crosswalk keys
<code>21-connectivity-full.png</code>	Connectivity, entire page
<code>22-connectivity-fold.png</code>	Connectivity, above the fold
<code>23-episode-page-full.png</code>	Episode page, entire page
<code>24-episode-page-fold.png</code>	Episode page, above the fold
<code>25-explain-decide-idle.png</code>	Explain and Decide, idle
<code>26-episode-page-masthead.png</code>	Episode page masthead
<code>27-dashboard-full-expanded.png</code>	Operations with every accordion open

Appendix B — The demo claim

E-000013 is the strongest single case in the dataset and every screen in this document uses it.

BEVACIZUMAB 400 MG/16 ML · 400 units · medical benefit · Blue Harbor Health · dispensed 20 Aug 2025
Insurance: **B-10** — denied by the payer, no appeal filed. \$0.00 received of \$39,360.00 expected.
340B rebate: **C-08** — approved, paid, cash matched. \$15,000.00 received of \$15,000.00 expected.
Flag **X-1** . Reopened from **PENDING** . 8 records · 2 cash movements · 4 verdict changes across 11 evaluations.

The insurer refused to pay for the drug. The manufacturer sent the 340B rebate anyway. The hospital is holding \$15,000 it may owe back, on a claim that paid it nothing — and is still owed \$39,360 it will probably never see.

No single-track system can see this. That is the product in one claim.

Do not use E-000042 or the A-07 / C-09 pair in any demo. Both walkthrough documents open with a correction banner: those figures were hand-composed and never re-derived from a run, and that pair lands on no episode in the database. Episode ids are reassigned on every reseed, so any id quoted in prose is a snapshot of one dataset generation. The figures above were read off the running app on 22 September 2026 and carry the same caveat.